

The ARM logo is displayed in white lowercase letters. The background of the slide is a stylized, isometric illustration of a city at night, with buildings and streets rendered in shades of blue and teal, and numerous small orange and yellow dots representing city lights.

arm

Debugging and Optimization Tools for Python HPC Codes

Beau Paisley<Beau.Paisley@arm.com>

RMACC 2021

Arm Forge == DDT + MAP + Performance Reports

An interoperable toolkit for debugging and performance engineering



Commercially supported
by Arm



Fully Scalable



Very user-friendly

The de-facto standard for HPC development

- Most widely-used debugging and profiling suite in HPC
- Integration of debugger DDT, source code profiler MAP, and application profiler Performance Reports
- Fully supported by Arm on Intel, AMD, Arm, IBM Power, Nvidia GPUs, etc.
- Multithreaded and multi-process support for C/C++, Fortran and Python

State-of-the art debugging and profiling capabilities

- Powerful and in-depth error detection mechanisms (including memory debugging)
- Sampling-based profiler to identify and understand bottlenecks
- Available at any scale

Easy to use by everyone

- Unique capabilities to simplify remote interactive sessions
- Innovative approach to present quintessential information to users

The Scalable DDT GUI

The screenshot displays the Arm DDT - Arm Forge 20.2 interface. At the top, a menu bar includes File, Edit, View, Control, Tools, Window, and Help. Below it is a toolbar with various icons for running, stepping, and debugging. A status bar at the top indicates the current group is 'All' and provides options for focus (Group, Process, Thread) and step-through settings.

On the left, a 'Project Files' pane shows a tree view with 'Application Code', 'Sources', and 'External Code'. The 'Sources' folder is expanded, showing 'heightmap.c'. A search bar (Ctrl+K) is also present.

The central pane displays the source code of 'heightmap.c'. A green callout box labeled 'Source code & breakpoints' points to the code. The code includes comments and function calls like 'calculate', 'malloc', 'memset', and 'free'. A red circle indicates a breakpoint set at line 77.

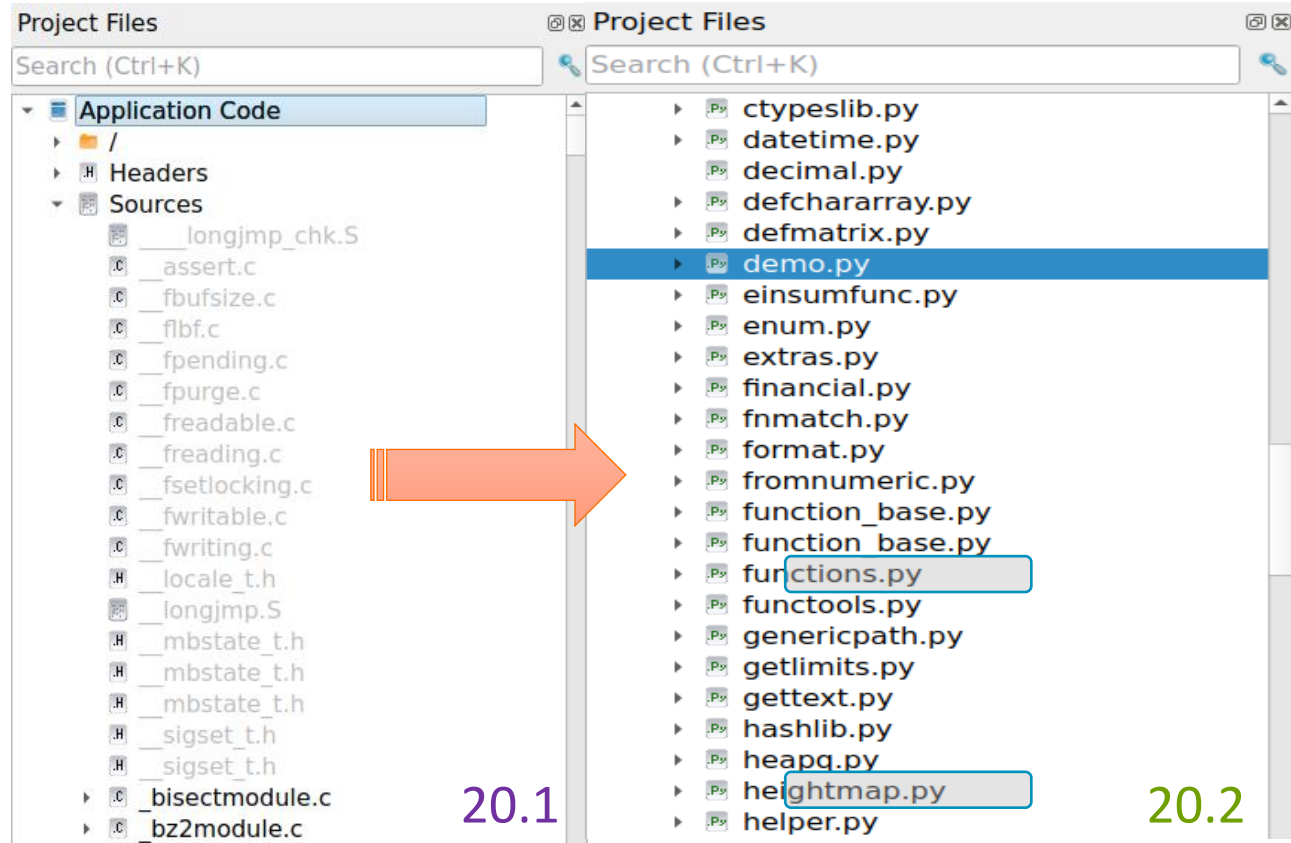
On the right, a 'Locals' pane shows the 'Current Line(s)' and 'Current Stack'. The 'Current Line(s)' pane lists variables: 'heightmap' (0x732a90), 'y' (1), 'blockSize' (30), 'x' (0), 'xOffset' (30), and 'yOffset' (0). An orange button labeled 'Variables' and a blue button labeled 'Sparklines' are located below this pane.

At the bottom, a 'Stacks' pane shows the current stack frame: 'main (heightmap.c:77)' and 'wave (heightmap.c:22)'. A blue callout box labeled 'Aggregate data across processes / threads' points to this pane.

A yellow callout box at the top right says 'Inspecting & controlling multiple processes'.

The bottom right corner shows a 'Ready' status.

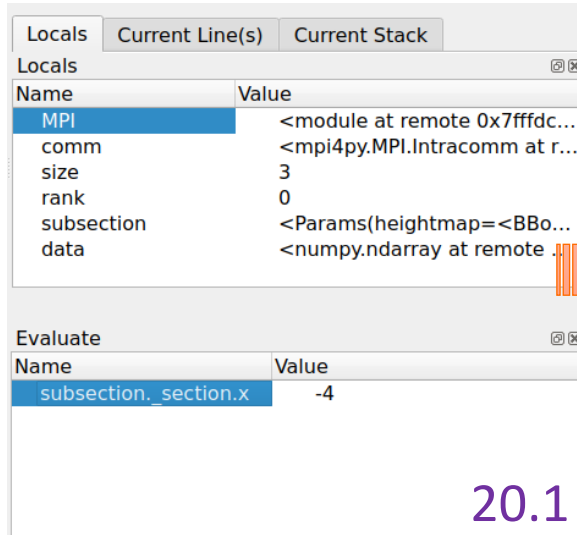
Project Files



```
demo.py x
1 import argparse # To parse --mpi option ...
2 import heightmap # BBox, Params and calculate
3 import functions # Predefined functions to calculate t
4 import numpy as np
5
6 params = heightmap.Params()
7 params.heightmap.size(9, 9).center()
8
9 def main_mpi():
10     """Splits the sections evenly across multiple proc
11     from mpi4py import MPI
12     print("*** MPI ***")
13     global params
14     comm = MPI.COMM_WORLD
15     size = comm.Get_size()
16     rank = comm.Get_rank()
17     subsection = params.subsection(size, rank)
18     data = heightmap.calculate(subsection, func)
19     print(data)
```

Inspecting Objects

New Python debugging features in DDT 20.2



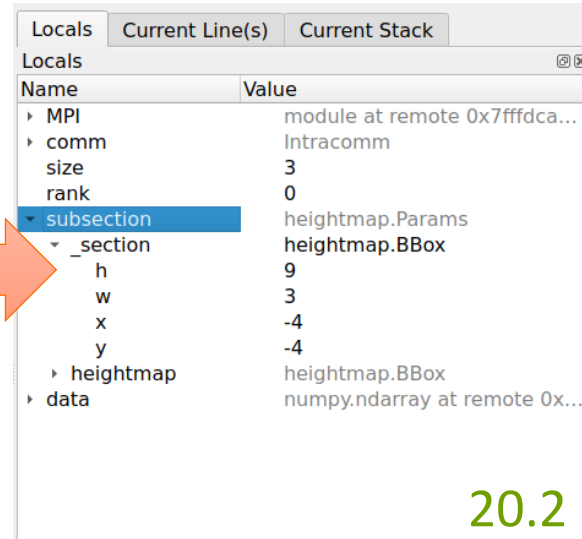
Locals

Name	Value
MPI	<module at remote 0x7ffdc...
comm	<mpi4py.MPI.Intracomm at r...
size	3
rank	0
subsection	<Params(heightmap=<BBo...
data	<numpy.ndarray at remote...

Evaluate

Name	Value
subsection._section.x	-4

20.1



Locals

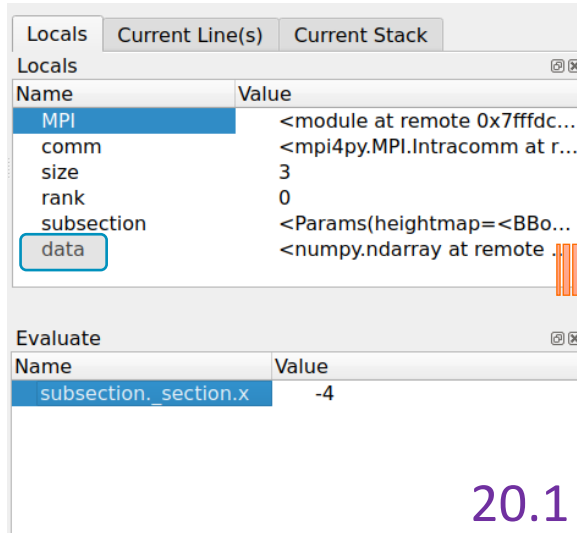
Name	Value
▶ MPI	module at remote 0x7ffdc...
▶ comm	Intracomm
size	3
rank	0
▼ subsection	heightmap.Params
▼ _section	heightmap.BBox
h	9
w	3
x	-4
y	-4
▶ heightmap	heightmap.BBox
▶ data	numpy.ndarray at remote 0x...

20.2

```
demo.py x
1 import argparse # To parse --mpi option ...
2 import heightmap # BBox, Params and calculate
3 import functions # Predefined functions to calculate t!
4 import numpy as np
5
6 params = heightmap.Params()
7 params.heightmap.size(9, 9).center()
8
9 def main_mpi():
10     """Splits the sections evenly across multiple proc
11     from mpi4py import MPI
12     print("*** MPI ***")
13     global params
14     comm = MPI.COMM_WORLD
15     size = comm.Get_size()
16     rank = comm.Get_rank()
17     subsection = params.subsection(size, rank)
18     data = heightmap.calculate(subsection, func)
19     print(data)
```

Inspecting Modules

New Python debugging features in DDT 20.2



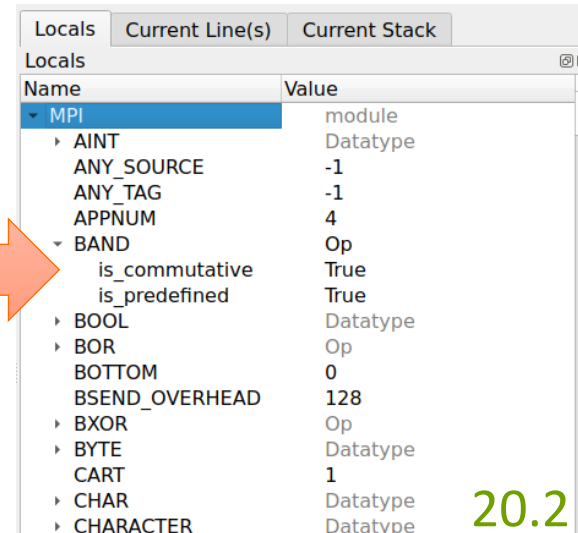
Locals

Name	Value
MPI	<module at remote 0x7ffdc...
comm	<mpi4py.MPI.Intracomm at r...
size	3
rank	0
subsection	<Params(heightmap=<BBo...
data	<numpy.ndarray at remote...

Evaluate

Name	Value
subsection._section.x	-4

20.1



Locals

Name	Value
MPI	module
▶ AINT	Datatype
▶ ANY_SOURCE	-1
▶ ANY_TAG	-1
▶ APPNUM	4
▶ BAND	Op
is_commutative	True
is_predefined	True
▶ BOOL	Datatype
▶ BOR	Op
▶ BOTTOM	0
▶ BSEND_OVERHEAD	128
▶ BXOR	Op
▶ BYTE	Datatype
▶ CART	1
▶ CHAR	Datatype
▶ CHARACTER	Datatype

20.2

```
demo.py x
1  import argparse # To parse --mpi option ...
2  import heightmap # BBox, Params and calculate
3  import functions # Predefined functions to calculate t
4  import numpy as np
5
6  params = heightmap.Params()
7  params.heightmap.size(9, 9).center()
8
9  def main_mpi():
10     """Splits the sections evenly across multiple proc
11     from mpi4py import MPI
12     print("*** MPI ***")
13     global params
14     comm = MPI.COMM_WORLD
15     size = comm.Get_size()
16     rank = comm.Get_rank()
17     subsection = params.subsection(size, rank)
18     data = heightmap.calculate(subsection, func)
19     print(data)
```

Inspecting NumPy Arrays

New Python debugging features in DDT 20.2

Locals	Current Line(s)	Current Stack
Locals		
Name	Value	
▶ MPI	module	
▶ comm	Intracomm	
size	3	
rank	0	
▶ subsection	heightmap.Params	
▼ data	ndarray of length 9	
▼ [0]	ndarray of length 3	
[0]	5.656854249492381	
[1]	5.0	
[2]	4.47213595499958	
▼ [1]	ndarray of length 3	
[0]	5.0	
[1]	4.242640687119285	
[2]	3.605551275463989	
▶ [2]	ndarray of length 3	
▶ [3]	ndarray of length 3	
▶ [4]	ndarray of length 3	

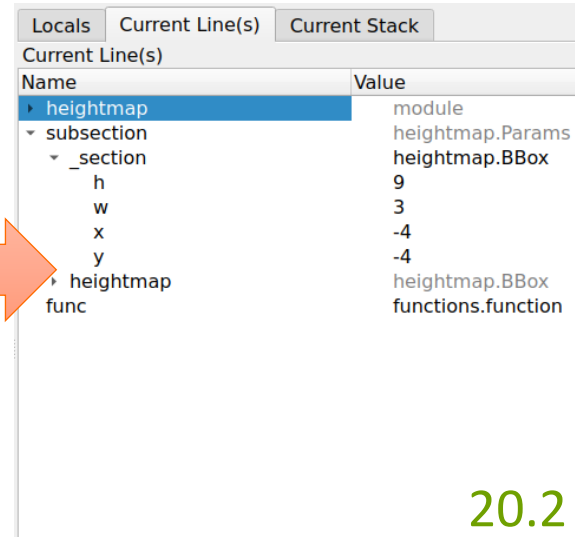

```
demo.py x
1  import argparse # To parse --mpi option ...
2  import heightmap # BBox, Params and calculate
3  import functions # Predefined functions to calculate t!
4  import numpy as np
5
6  params = heightmap.Params()
7  params.heightmap.size(9, 9).center()
8
9  def main_mpi():
10     """Splits the sections evenly across multiple proc
11     from mpi4py import MPI
12     print("*** MPI ***")
13     global params
14     comm = MPI.COMM_WORLD
15     size = comm.Get_size()
16     rank = comm.Get_rank()
17     subsection = params.subsection(size, rank)
18     data = heightmap.calculate(subsection, func)
19     print(data)
```



Not limited to NumPy arrays → Python sequence protocol!

Variables on Current Line(s)

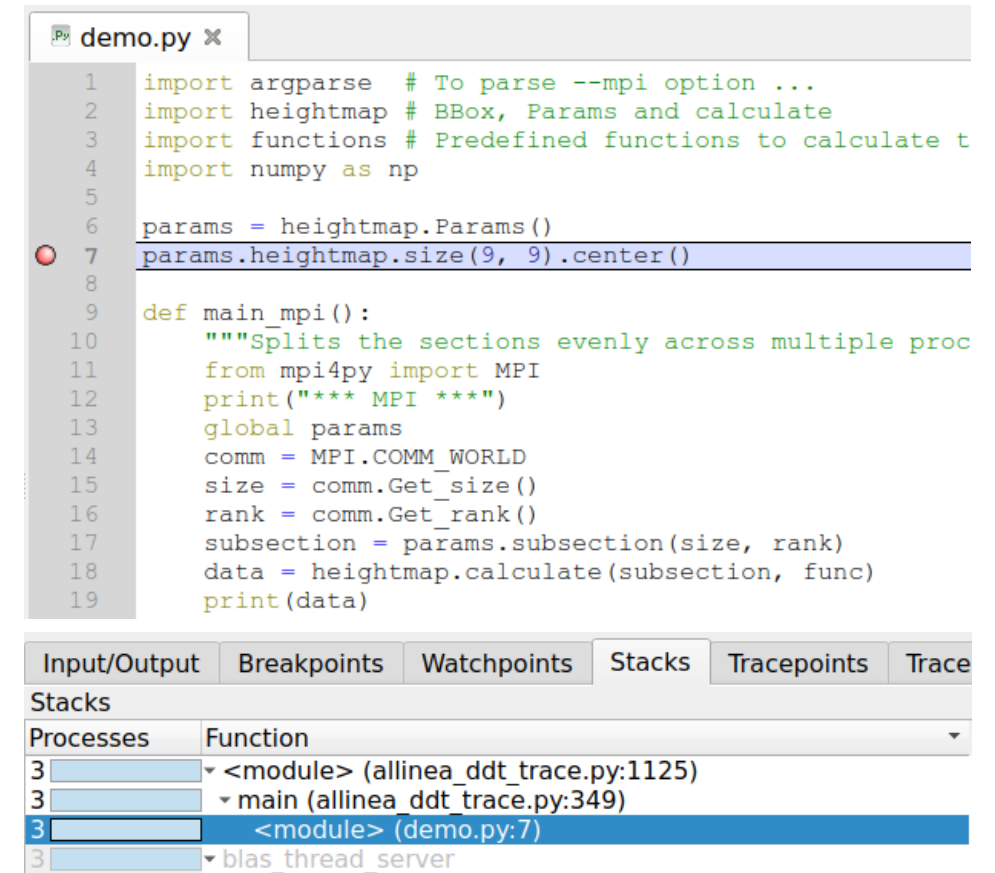
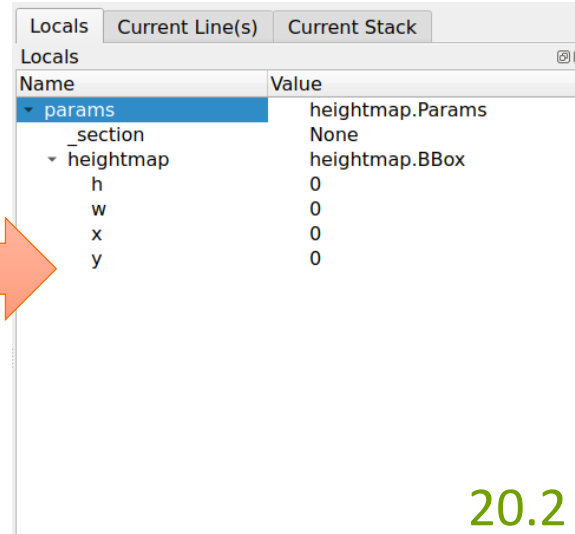

New Python debugging features in DDT 20.2



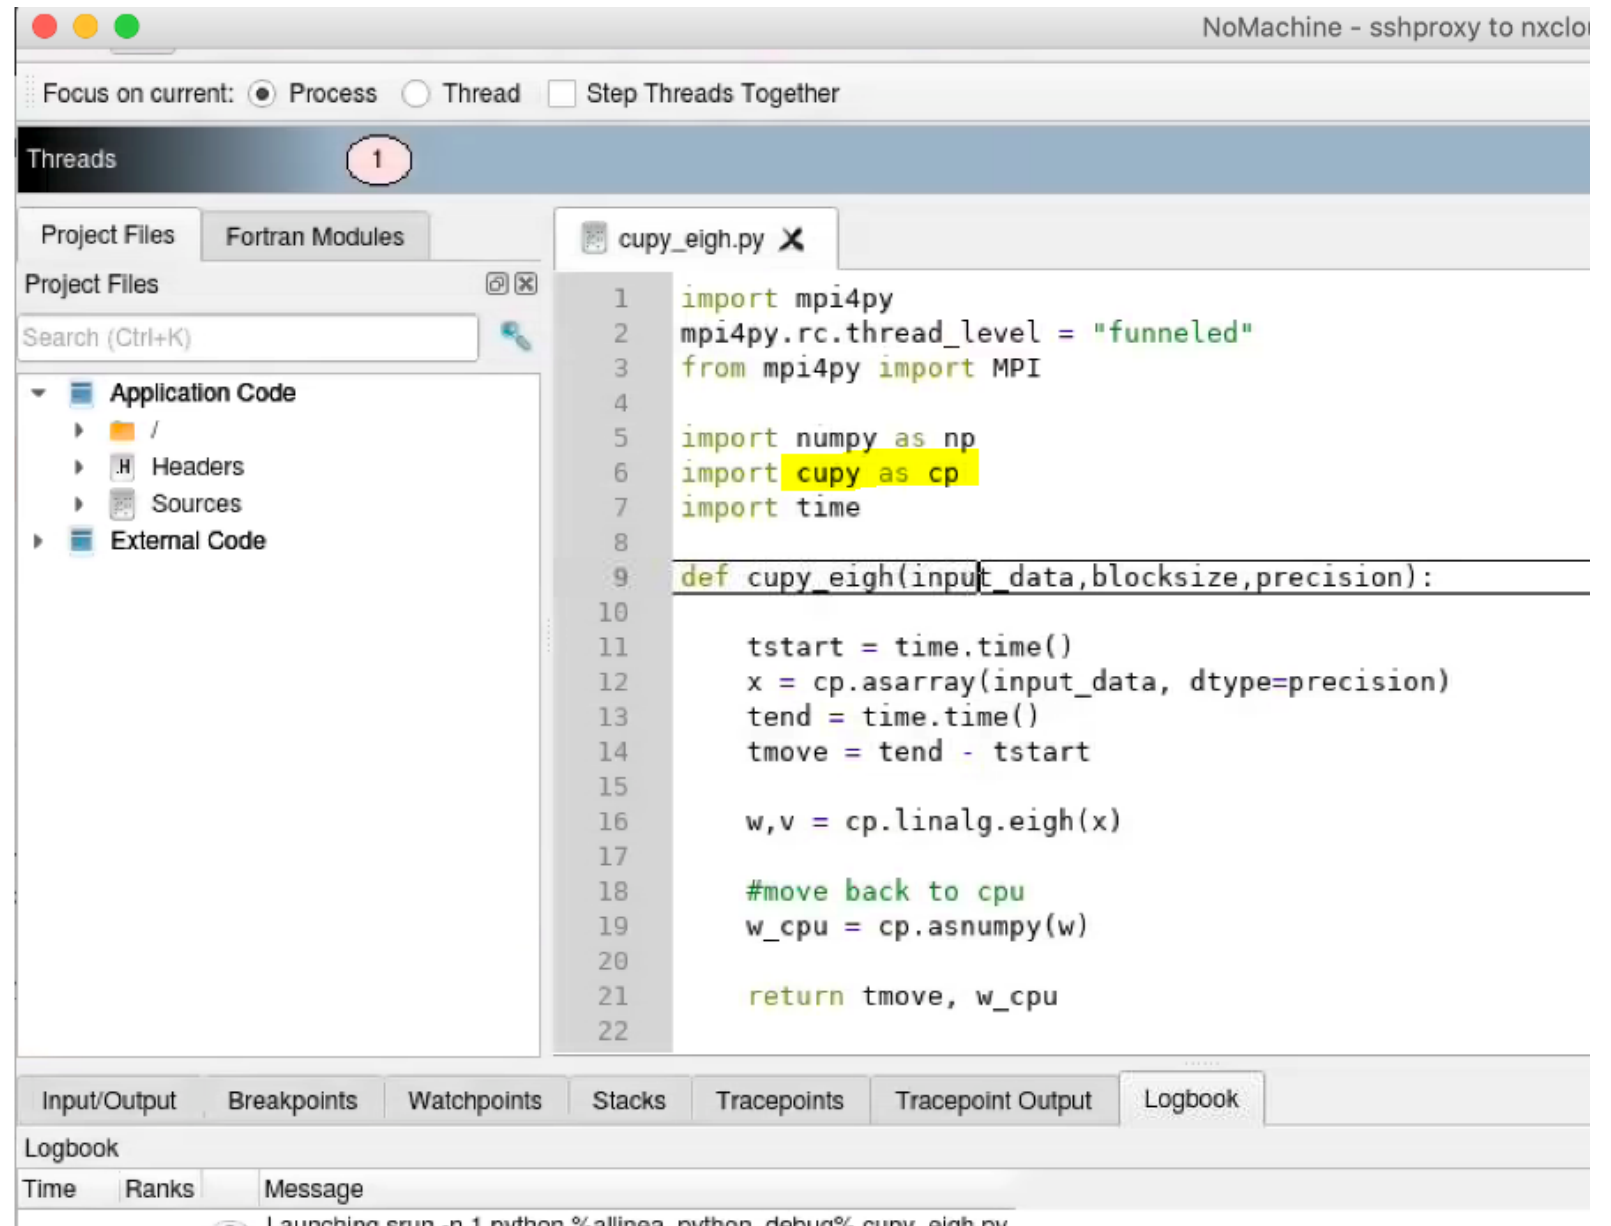
```
demo.py x
1 import argparse # To parse --mpi option ...
2 import heightmap # BBox, Params and calculate
3 import functions # Predefined functions to calculate t
4 import numpy as np
5
6 params = heightmap.Params()
7 params.heightmap.size(9, 9).center()
8
9 def main_mpi():
10     """Splits the sections evenly across multiple proc
11     from mpi4py import MPI
12     print("*** MPI ***")
13     global params
14     comm = MPI.COMM_WORLD
15     size = comm.Get_size()
16     rank = comm.Get_rank()
17     subsection = params.subsection(size, rank)
18     data = heightmap.calculate(subsection, func)
19     print(data)
```


Quick Script: Globals?

New Python debugging features in DDT 20.2



Debugging Python that Uses cuda-based Toolkits



9 Step guide: optimizing high performance applications

arm

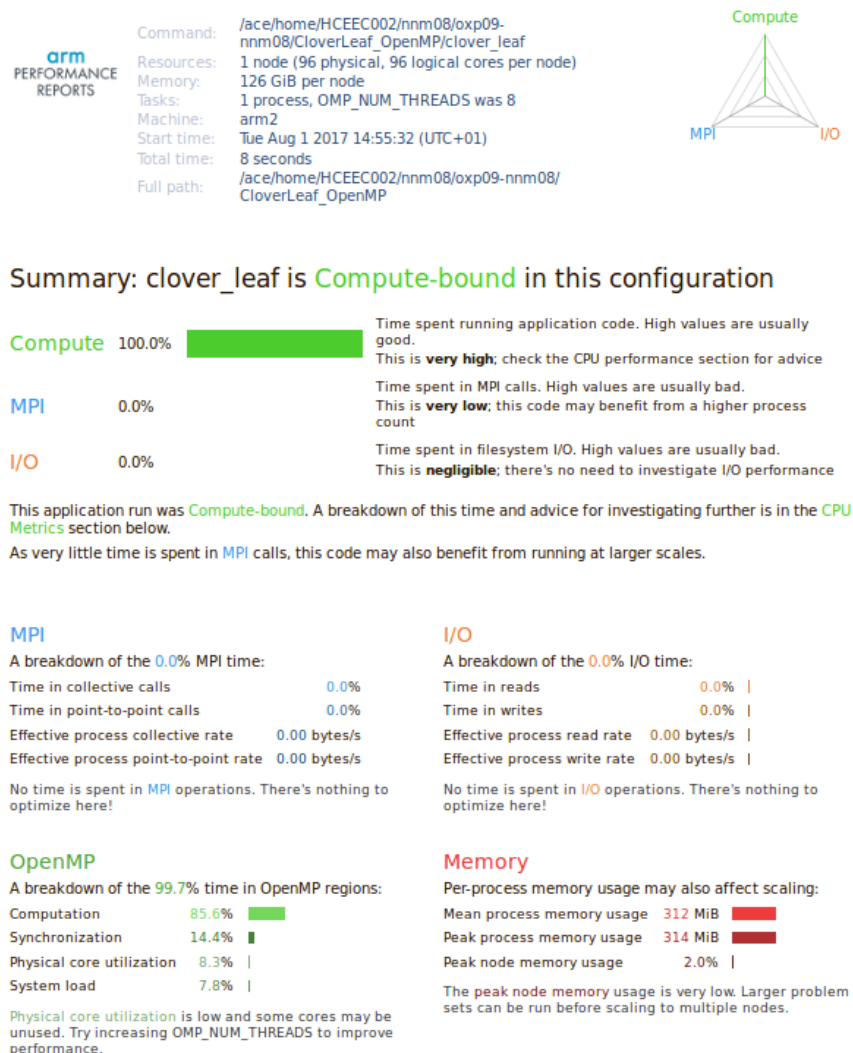
Improving the efficiency of your parallel software holds the key to solving more complex research problems faster. This pragmatic, 9 Step best practice guide will help you identify and focus on application readiness, bottlenecks and optimizations one step at a time.



Key:

✓ arm PERFORMANCE REPORTS
✓ arm FORGE

Performance Snapshots with Arm Performance Reports



No source code needed

Less than 5% runtime overhead


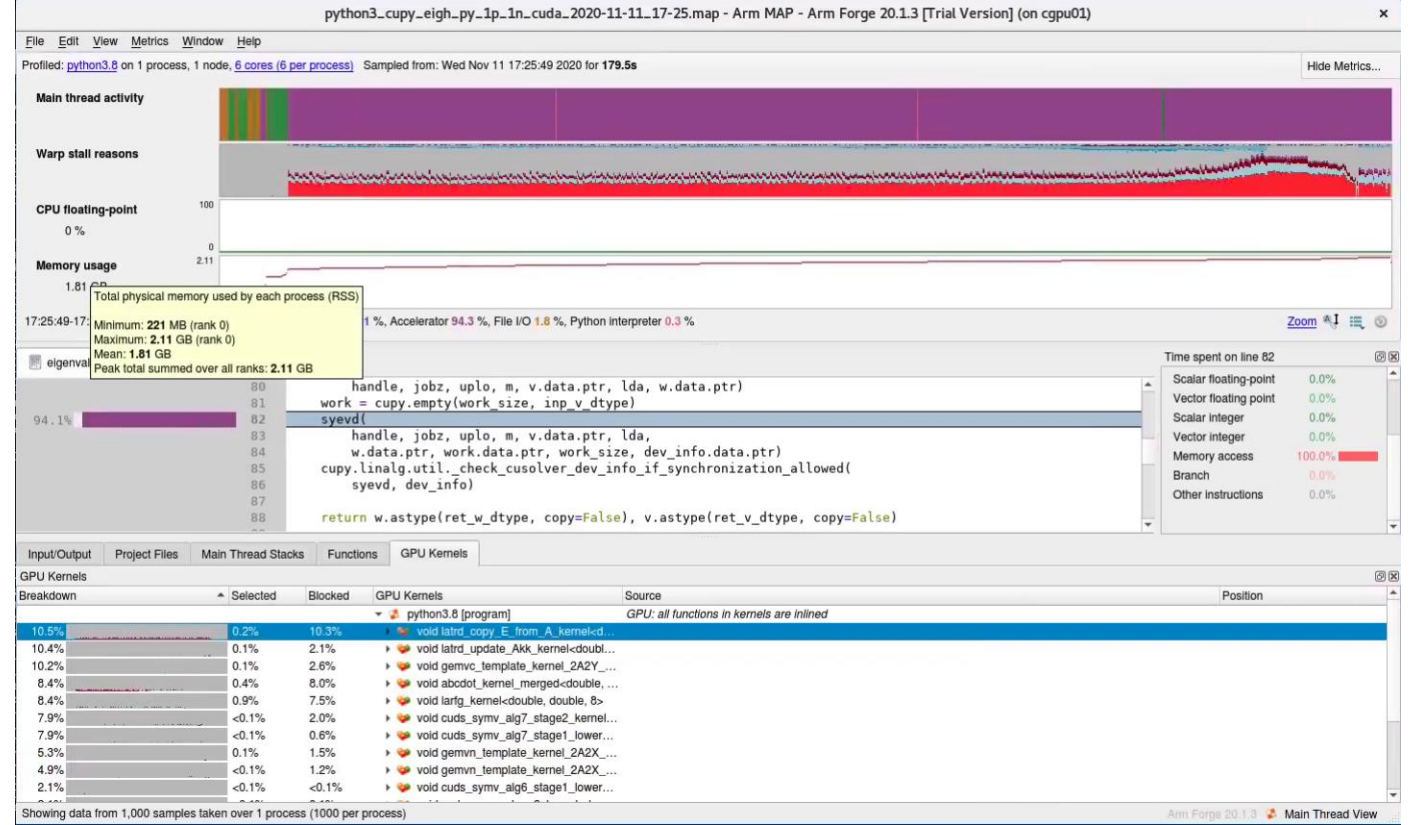
Fully scalable

Run regularly – or in regression tests

Explicit and usable output

MAP Highlights

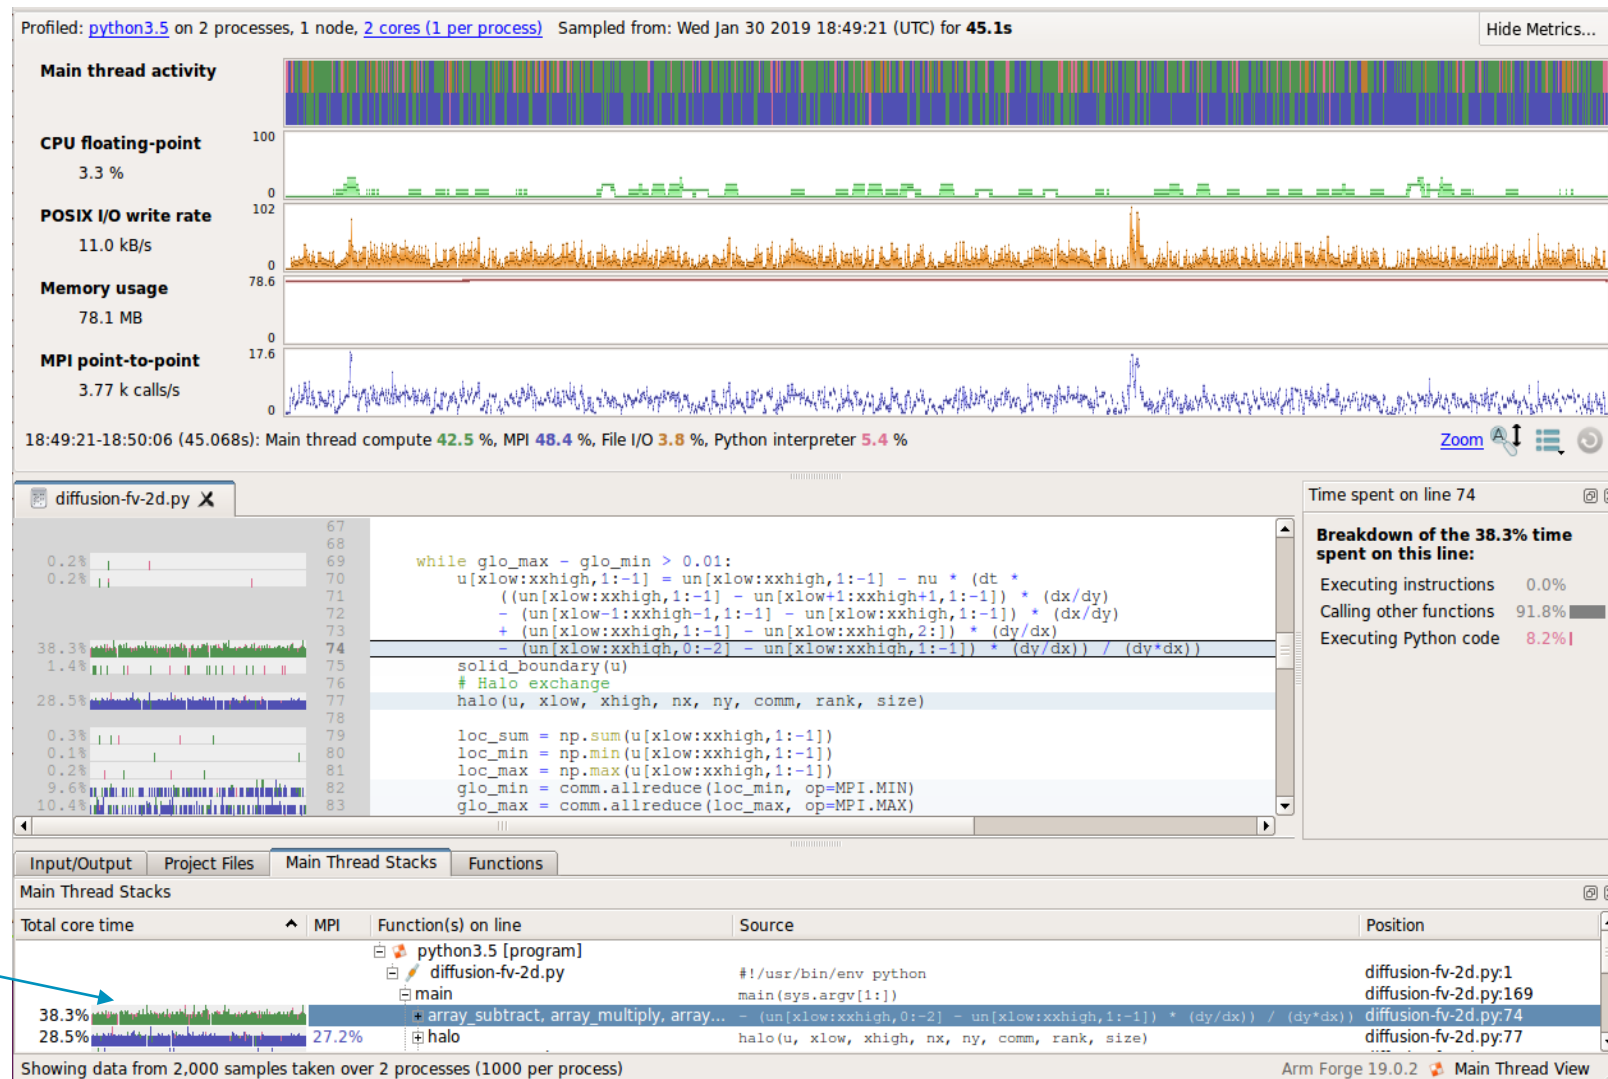
- MAP is a sampling based scalable profiler
 - Built on same framework as DDT
 - Parallel support for MPI, OpenMP, CUDA
 - Designed for C/C++/Fortran, and Python
- Designed for 'hot-spot' analysis
 - Stack traces
 - Augmented with performance metrics
- Adaptive sampling rate
 - Periodically resamples and modifies sampling frequency
 - Low overhead, scalable and small file size



Python Source Code Profiling with MAP

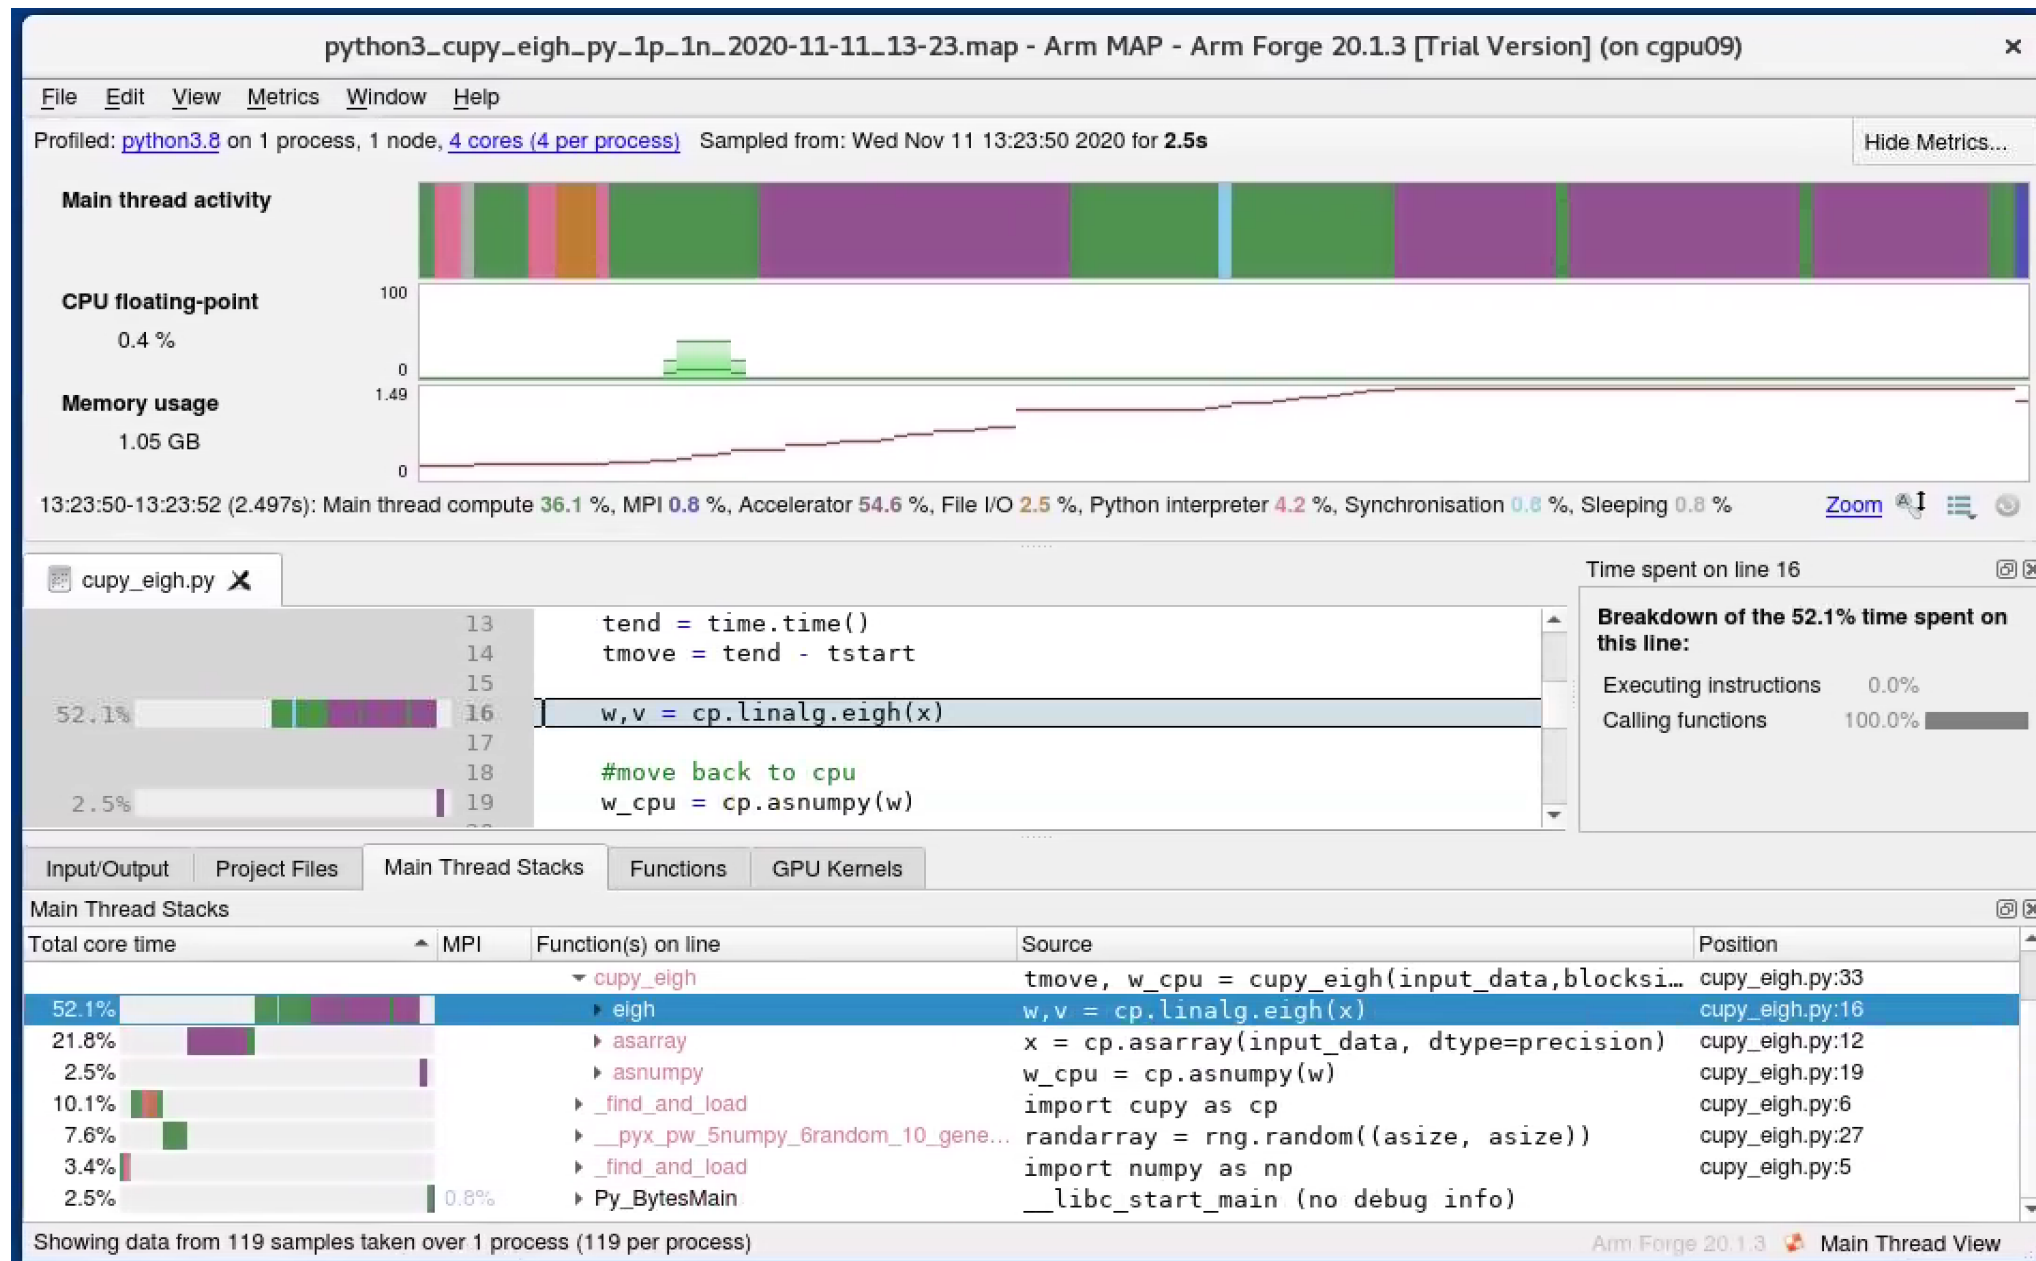
- Profile Pure and Mixed code
 - Call stacks
 - Time in interpreter
- Works with MPI4PY
 - Usual MAP metrics
- Source code view
 - Mixed language support

Note: Green as operation is on numpy array, so backed by C routine, not Python (which would be pink)




map --profile jsrun -n 2 python3 ./diffusion-fv-2d.py

Profiling Python that Uses Cuda-Based Toolkits



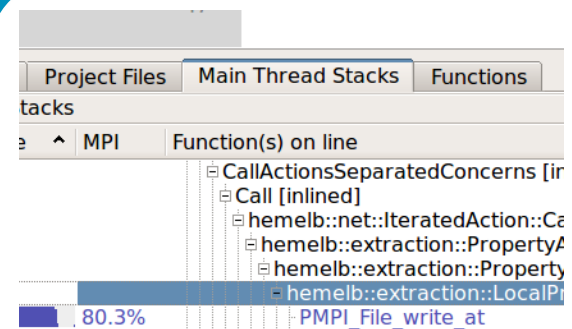
MAP Source Code Profiler Highlights




Find the peak memory use

```
30      ! late to the party
31      do j=1,20*nprocs; a
32      end if
33
34      if (pe /= 0) then
35      call MPI_SEND(a, si
36      else
37      do from=1,nprocs-1
38      call MPI_RECV(b,
39      do j=1,50; b=sqrt
40      print *, "Answer f
41      end do
42      end if
43      end do
44      call MPI_BARRIER(MPI CO
```


Fix an MPI imbalance



Remove I/O bottleneck



Make sure OpenMP regions make sense



Improve memory access

```
size, nproc, mat_a
A[i*size+k]*B[k*s

nalize();
/(size, mat c, file
```

Restructure for vectorization

Next Steps

- Contact the Arm HPC Tools team to learn more
- Schedule a live demo to see the tools in action
- Obtain a trial to have a look yourself (Most RMACC members have licenses for Forge)



The Arm trademarks featured in this presentation are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. All other marks featured may be trademarks of their respective owners.

www.arm.com/company/policies/trademarks

Beau Paisley, Beau.Paisley@arm.com, Solutions Architect, North America – HPC Tools

Andrew Westergren, Andrew.Westergren@arm.com Senior Manager, North America – HPC Tools